

Andrej Karpathy Research Report

Detailed Research Report Generated From Multiple Sources

Topic	Andrej Karpathy
Generated On	13 June 2026
Articles	0
News Items	10
Images	1

Contents

1. Executive Summary
2. Background Information
3. Timeline
4. Key Facts
5. Recent Developments
6. Analysis
7. References


Executive Summary

Andrej Karpathy (born 23 October 1986) is a Slovak-Canadian AI researcher, who co-founded and formerly worked at OpenAI, where he specialized in deep learning and computer vision. He also worked as the director of artificial intelligence and Autopilot Vision at Tesla, and in 2024 he founded Eureka Labs, an AI education platform. In 2026 he joined Anthropic as part of the pretraining team.

Key Topics

- Karpathy
- Anthropic
- Openai
- Team
- Tesla
- Where
- Andrej
- What
- Coding
- Will
- Research
- Vibe
- Code
- Model
- Agent

Image Gallery



Background Information

Education and early life

Karpathy was born in Bratislava, Czechoslovakia (now Slovakia), and moved with his family to Toronto when he was 15. He completed his Computer Science and Physics bachelor's degrees at University of Toronto in 2009 and his master's degree at University of British Columbia in 2011, where he worked on physically simulated figures (for example, a simulated runner or a simulated person in a crowd) with his adviser Michiel van de Panne. In 2006, Karpathy began posting videos on YouTube on his channel, badmephisto. He garnered fame by posting Rubik's cube tutorials which have been used by famous speedcubers such as Feliks Zemdegs. The channel has over 9 million views as of June 2025. Karpathy received a PhD from Stanford University in 2015 under the supervision of Fei-Fei Li, focusing on the intersection of natural language processing and computer vision, and deep learning models suited for this task.

Career and research

He authored and was the primary instructor of the first deep learning course at Stanford, CS 231n: Convolutional Neural Networks for Visual Recognition. The course became one of the largest classes at Stanford, growing from 150 students in 2015 to 750 in 2017. Karpathy is a founding member of the artificial intelligence research group OpenAI, where he was a research scientist from 2015 to 2017. In June 2017 he became Tesla's director of artificial intelligence and reported to Elon Musk. He was named one of MIT Technology Review's Innovators Under 35 for 2020. After taking a several-months-long sabbatical from Tesla, he announced he was leaving the company in July 2022. As of February 2023, he makes YouTube videos on how to create artificial neural networks. On February 9, 2023, Karpathy announced he was returning to OpenAI. A year later on February 13, 2024, an OpenAI spokesperson confirmed that Karpathy had left OpenAI. In the same year, he was named one of Time Magazine's 100 Most Influential People in AI. On July 16, 2024, Karpathy announced on his X account that he started a new AI education company called Eureka Labs. Their first product was the AI course, LLM101n. He also has a broader educational effort, the "Zero to Hero" series on LLM fundamentals. The company also advocates for AI teaching assistants, a concept which has been criticized due to data privacy concerns and the removal of personal connection between teacher and student. In February 2025, Karpathy coined the term vibe coding to describe a way of building software where you describe what you want in plain language, let an LLM generate the code, and iterate from there rather than writing code by hand. The term moved fast: Merriam-Webster added it as a "slang & trending" expression in March 2025, and Collins named it Word of the Year for 2025. The idea built on something Karpathy had said in 2023 -- that "the hottest new programming language is English" -- suggesting developers would increasingly direct computers in natural language rather than code. By early 2025, Y Combinator reported that 25% of startups in its Winter batch had codebases that were 95% AI-generated. On May 19, 2026, he announced that he joined Anthropic via a statement on X, while the company stated that he will be leading a team for research in pretraining.

External links

Official website of Eureka Labs Official GitHub account of Eureka Labs Andrej Karpathy at GitHub

Timeline

Year	Mentions
1986	1
2006	1
2009	1
2011	1
2015	5
2017	6
2020	1
2022	4
2023	5
2024	7
2025	7
2026	5

Key Facts

- Andrej Karpathy (born 23 October 1986) is a Slovak-Canadian AI researcher, who co-founded and formerly worked at OpenAI, where he specialized in deep learning and computer vision.
- He also worked as the director of artificial intelligence and Autopilot Vision at Tesla, and in 2024 he founded Eureka Labs, an AI education platform.
- In 2026 he joined Anthropic as part of the pretraining team.
- Karpathy was born in Bratislava, Czechoslovakia (now Slovakia), and moved with his family to Toronto when he was 15.
- He completed his Computer Science and Physics bachelor's degrees at University of Toronto in 2009 and his master's degree at University of British Columbia in 2011, where he worked on physically simulated figures (for example, a simulated runner or a simulated person in a crowd) with his adviser Michiel van de Panne.
- In 2006, Karpathy began posting videos on YouTube on his channel, badmephisto.
- He garnered fame by posting Rubik's cube tutorials which have been used by famous speedcubers such as Feliks Zemdegs.
- Karpathy received a PhD from Stanford University in 2015 under the supervision of Fei-Fei Li, focusing on the intersection of natural language processing and computer vision, and deep learning models suited for this task.
- He authored and was the primary instructor of the first deep learning course at Stanford, CS 231n: Convolutional Neural Networks for Visual Recognition.
- The course became one of the largest classes at Stanford, growing from 150 students in 2015 to 750 in 2017.
- Karpathy is a founding member of the artificial intelligence research group OpenAI, where he was a research scientist from 2015 to 2017.
- In June 2017 he became Tesla's director of artificial intelligence and reported to Elon Musk.
- He was named one of MIT Technology Review's Innovators Under 35 for 2020.
- After taking a several-months-long sabbatical from Tesla, he announced he was leaving the company in July 2022.
- As of February 2023, he makes YouTube videos on how to create artificial neural networks.
- On February 9, 2023, Karpathy announced he was returning to OpenAI.
- A year later on February 13, 2024, an OpenAI spokesperson confirmed that Karpathy had left OpenAI.
- In the same year, he was named one of Time Magazine's 100 Most Influential People in AI.
- On July 16, 2024, Karpathy announced on his X account that he started a new AI education company called Eureka Labs.
- He also has a broader educational effort, the "Zero to Hero" series on LLM fundamentals.
- The company also advocates for AI teaching assistants, a concept which has been criticized due to data privacy concerns and the removal of personal connection between teacher and student.

- In February 2025, Karpathy coined the term vibe coding to describe a way of building software where you describe what you want in plain language, let an LLM generate the code, and iterate from there rather than writing code by hand.
- The term moved fast: Merriam-Webster added it as a "slang & trending" expression in March 2025, and Collins named it Word of the Year for 2025.
- The idea built on something Karpathy had said in 2023 -- that "the hottest new programming language is English" -- suggesting developers would increasingly direct computers in natural language rather than code.
- By early 2025, Y Combinator reported that 25% of startups in its Winter batch had codebases that were 95% AI-generated.
- On May 19, 2026, he announced that he joined Anthropic via a statement on X, while the company stated that he will be leading a team for research in pretraining.
- Official website of Eureka Labs Official GitHub account of Eureka Labs Andrej Karpathy at GitHub Andrej Karpathy, who coined "vibe coding," now deems it obsolete, advocating for "agentic engineering" after joining Anthropic.
- Vibe coding allowed AI to build code from plain English, accelerating prototypes.
- Agentic engineering elevates this, with AI writing, testing, debugging, and deploying code autonomously, requiring human direction.
- This paradigm shift dramatically shortens development cycles, enabling individuals to achieve more and reducing the cost of innovation.
- It's already transforming industries, including regulated banking.
- Founders are encouraged to embrace agentic engineering by directing AI agents for projects, learning from their limitations, and prioritizing the underlying AI model's quality.
- This new era emphasizes shipping complex solutions without needing full human understanding.
- Andrej Karpathy gave vibe coding its name in February 2025 .
- In May he joined Anthropic to rebuild its pretraining research team from the inside.
- The man who coined the term has already named the next phase.
- If you have been telling clients you vibe code, the phrasing might be a year out of date.
- Vibe coding meant typing what you wanted in plain English and letting the model build it.
- It worked because the first wave of AI coding tools were fast at scaffolding and forgiving on small mistakes.
- You could ship a landing page in an afternoon and a prototype in a weekend.

Recent Developments

1. Is Vibe Coding Already Dead? Even Karpathy Is Moving On - Forbes

Forbes | 2026-06-12

Andrej Karpathy, who coined "vibe coding," now deems it obsolete, advocating for "agentic engineering" after joining Anthropic. Vibe coding allowed AI to build code from plain English, accelerating prototypes. Agentic engineering elevates this, with AI writing, testing, debugging, and deploying code autonomously, requiring human direction. This paradigm shift dramatically shortens development cycles, enabling individuals to achieve more and reducing the cost of innovation. It's already transforming industries, including regulated banking. Founders are encouraged to embrace agentic engineering by directing AI agents for projects, learning from their limitations, and prioritizing the underlying AI model's quality. This new era emphasizes shipping complex solutions without needing full human understanding. Andrej Karpathy gave vibe coding its name in February 2025. By February 2026 he had declared it outdated and moved on. In May he joined Anthropic to rebuild its pretraining research team from the inside. The man who coined the term has already named the next phase. He calls it agentic engineering. If you have been telling clients you vibe code, the phrasing might be a year out of date. Vibe coding meant typing what you wanted in plain English and letting the model build it. It worked because the first wave of AI coding tools were fast at scaffolding and forgiving on small mistakes. You could ship a landing page in an afternoon and a prototype in a weekend. Agentic engineering is the next layer up. The model writes the code, runs the tests, reads the errors, fixes the code, runs the tests again, and reports back when something is shipped. Your role becomes direction. Karpathy's argument is that vibe coding was a bridge phase. The end state was always going to be a system where the human sets the goal and the agent does the work. Bloomberg ran a video in early May where senior engineers at three different companies agreed the change had already happened inside their teams. None of them used the words vibe coding to describe how they ship today. Harvard Gazette covered the same story from the academic side. Researchers there are using agentic engineering to build small internal tools that would have needed a developer to start six months ago. The people running these experiments are scientists, not coders. They write a brief, hand it to an agent, and pick up the working tool an hour later. Founders adopted vibe coding for the same reason they adopted Canva. They wanted to ship faster and without hiring. The tool did the job well enough to feel like a superpower. Agentic engineering raises the ceiling on what one person can build. You ship working features now. Deployed, with tests, that a paying customer can use. That is a different roadmap. It changes what you can promise. It changes what one person can charge for. The first thing that changes when you start running agents is cycle time. Building a feature drops from a week to a day, and your roadmap doubles in scope without doubling the team. The cost of trying a new idea drops to almost nothing in the same move. You can spin up an agent to test a feature and pivot before lunch. Knowing how to break a project into clean tasks and review the output is worth more than knowing how to write the code yourself. PYMNTS reported in May that vibe coding had already crossed into regulated banking, with one US neobank shipping a customer-facing feature built entirely by an agent. The compliance team approved it after a four-hour review. A year ago that build would have taken four months. The regulators are catching up to a model that already exists in production. Pick one project on your roadmap that has been stuck. The one you keep meaning to ship and keep deprioritising because it would take a week you do not have. Open Claude Code, Cursor, or another agent-based tool. Write the spec as if you were briefing a freelancer. One paragraph on what it should do. Three bullet points on what good looks like. One sentence on what it should not do. Run the agent. Read what it produces. Send it back with one correction. Run it again. You will hit the limits within an hour. That is the point. You learn agentic

engineering by directing one and watching where it gets stuck. The patterns you see in those failures are what your next prompt fixes. The skill compounds quickly once you stop treating the agent like a search bar. Once you've learned from your first round, set the target on what you'll build next. Track which categories of task the agent finishes on the first attempt, which need one round of corrections, and which need you to take over. That spreadsheet is your map of where agentic engineering pays off in your business. The founders who run this loop fifty times this quarter will know things in October that other founders will not know until next year. Karpathy moved to Anthropic for a reason. The model is the bottleneck on what an agent can do. The companies building the best models are pulling ahead because their tool finishes more of a task without a human stepping in to correct. Pick your agent tool based on the model under it. The agent that finishes the most work on the first attempt is the one worth paying for. Change tools when the model under it stops winning. The vibe coding era taught you that you can build things you do not understand. The agentic engineering era will teach you something harder. You can ship things you do not understand. That is the real bet you are making this year. Be ready for it.

2. OpenAI co-founder Andrej Karpathy joins Anthropic's pre-training team - TechCrunch

TechCrunch | 2026-05-19

Andrej Karpathy, the AI researcher who co-founded and formerly worked at OpenAI and previously led AI at Tesla, has joined Anthropic. "I've joined Anthropic," Karpathy posted on X Tuesday. "I think the next few years at the frontier of LLMs will be especially formative. I am very excited to join the team here and get back to R&D." Personal update: I've joined Anthropic. I think the next few years at the frontier of LLMs will be especially formative. I am very excited to join the team here and get back to R&D. I remain deeply passionate about education and plan to resume my work on it in time. Karpathy started this week at Anthropic, where he is working on pre-training under team lead Nick Joseph. Pre-training is responsible for the large-scale training runs that give Claude its core knowledge and capabilities, according to the company. It's also one of the most expensive, compute-intensive phases of building a frontier model. An Anthropic spokesperson told TechCrunch that Karpathy will start a team focused on using Claude to accelerate pre-training research. Karpathy is one of the few researchers who can bridge the gap between LLM theory and large-scale training practice. Tapping him to build such a team is a clear sign from Anthropic that it believes AI-assisted research, rather than pure compute, is how it stays competitive with OpenAI and Google. While at OpenAI, Karpathy focused on deep learning and computer vision until he departed in 2017 to join Tesla. He led Tesla's Full Self-Driving (FSD) and Autopilot programs before leaving in 2022. He then went back to OpenAI for one year before leaving again in 2024 to start Eureka Labs, a startup dedicated to applying AI assistants to education. Karpathy hasn't shared many updates on Eureka Labs since its launch, and it's not clear if the renowned researcher will continue with the startup. He has also taught an online course called Neural Networks: Zero to Hero that helps students learn to build neural networks from scratch in code, and he has a YouTube channel where he semi-regularly posts lectures on LLMs and AI. "I remain deeply passionate about education and plan to resume my work on it in time," Karpathy said. Separately, Anthropic has also brought on Chris Rohlfs to its frontier red team, which stress-tests advanced AI models against severe threats. Rohlfs is a veteran of the cybersecurity industry with more than 20 years of experience. He previously worked at Yahoo's well-respected cybersecurity team known as "The Paranoids," and more recently at Meta, where he worked for six years before joining Anthropic. Rohlfs was also a fellow at Georgetown's Center for Security and Emerging Technology, where he worked on the CyberAI project. "We have a real opportunity in front of us to dramatically improve cyber security with AI," Rohlfs said in a

post on X . “I can’t think of a better company or team to join at this critical moment in time.” When you purchase through links in our articles, we may earn a small commission . This doesn’t affect our editorial independence. Rebecca Bellan is a senior reporter at TechCrunch where she covers the business, policy, and emerging trends shaping artificial intelligence. Her work has also appeared in Forbes, Bloomberg, The Atlantic, The Daily Beast, and other publications. You can contact or verify outreach from Rebecca by emailing rebecca.bellan@techcrunch.com or via encrypted message at [rebeccabellan.491](https://signal.me/#x/rebeccabellan.491) on Signal. Lorenzo Franceschi-Bicchierai is a Senior Writer at TechCrunch, where he covers hacking, cybersecurity, surveillance, and privacy. You can contact or verify outreach from Lorenzo by emailing lorenzo@techcrunch.com , via encrypted message at +1 917 257 1382 on Signal, and [@lorenzofb](https://keybase.io/lorenzofb) on Keybase/Telegram. Get an inside look at what it takes to scale and succeed from leaders at Mach Industries, Founders Fund, and Shinkei Systems. Through candid fireside chats and high-impact networking, you’ll walk away with valuable insights and new connections. Jeff Bezos’s Prometheus raises \$12B to build an ‘artificial general engineer’ for the physical world Cybersecurity researchers aren’t happy about the guardrails on Anthropic’s Fable WWDC 2026: Everything announced on Siri AI, iOS 27, Apple Intelligence, and more Anthropic’s Claude Fable 5 is a version of Mythos the public can access today Microsoft’s open source tools were hacked to steal passwords of AI developers

3. Andrej Karpathy, OpenAI founding member and inventor of ‘vibe coding,’ defects to Anthropic - Fortune

Fortune | 2026-05-19

Anthropic can’t seem to stop winning. After a string of blockbuster model releases, a new funding round reportedly in talks at a valuation approaching \$1 trillion, and an annual run rate that’s nearly parabolic, it has now hired one of OpenAI’s—its bitter competitor—most famous alumni. “Personal update: I’ve joined Anthropic,” Andrej Karpathy wrote on X on Tuesday, in a post that drew nearly 3 million views within one hour. He said the next few years at the frontier of LLMs would be “especially formative” and that he was eager to get back to research. He started this week on the pretraining team. The decision suggests that Karpathy, whose writing on AI is followed by nearly 2 million people on X, has decided to put his stake in the AI race. He was first a founding member of OpenAI in 2015, left to run AI at Tesla , came back in 2023, and left only a year later to start his own education company, Eureka Labs. Karpathy has been well-known in AI for a decade, but the thing—or phrase—that etched him into AI legend was a tweet from last year. In February 2025, he posted that there was a “new kind of coding” that he called vibe coding : Describe what you want plainly and let the model do the work. The phrase escaped the industry and infected the business world, which turned its back against software companies and raced to develop its own bespoke agents, touching off the much-debated “SaaSocalypse” in its wake. (What wasn’t debated is that tens of billions of dollars in stock valuations evaporated as firms tried to vibe code their own solutions.) Collins Dictionary named it Word of the Year. The model he cited in that original tweet was Anthropic’s. As an Anthropic employee, Karpathy will be building on the work from another viral post. In March, he wired up an AI coding agent, handed it a single small language model, and let it run unsupervised for two days, testing and tweaking the training code on its own. After 700 experiments and 20 self-found optimizations, he said the same tweaks applied to a larger model cut training time by 11%. This, he said, was called autoresearch: “Part code, part sci-fi, and a pinch of psychosis” he wrote with a smiley face. The method would come to be known as “the Karpathy Loop.” Teaching that method looks, more or less, like his new job. According to Anthropic, Karpathy will be starting a team focused on using Claude to accelerate pretraining research, the large-scale training runs that give Claude its core knowledge and capabilities. He sits on a team led by Nick Joseph. Long before

any of this, though, Karpathy was famous for something else: his Rubik's Cube skills. He ran a YouTube channel called "badmephisto," where a generation of competitive cubers learned to "speedcube" by seeing the cube as 26 individual "cubies" rather than 54 colored stickers. By sticking to the small structure, he could move the whole thing. He solved a Rubik's Cube in about 17 seconds. Certainly the puzzles for Karpathy got harder—neural nets and then language models—but the method never really changed. Get a small enough system fully under control, and you can move something much bigger.

4. OpenAI Founding Member Andrej Karpathy Takes Role at Anthropic - Bloomberg.com

Bloomberg.com | 2026-05-19

OpenAI Founding Member Andrej Karpathy Takes Role at Anthropic Bloomberg.com

5. OpenAI co-founder and former Tesla AI executive Karpathy joins Anthropic - Reuters

Reuters | 2026-05-19

OpenAI co-founder and former Tesla AI executive Karpathy joins Anthropic Reuters

6. OpenAI co-founder Andrej Karpathy joins Anthropic - Axios

Axios | 2026-05-19

OpenAI co-founder Andrej Karpathy joins Anthropic Axios

7. Anthropic hires OpenAI co-founder Andrej Karpathy, former Tesla AI leader - CNBC

CNBC | 2026-05-19

Andrej Karpathy, an artificial intelligence researcher who co-founded OpenAI before getting poached by Tesla, announced on Tuesday that he's joining Anthropic. "I think the next few years at the frontier of LLMs will be especially formative," Karpathy wrote in a post on X, referring to large language models. "I am very excited to join the team here and get back to R&D." Anthropic said Karpathy starts this week and will be building a team focused on using Claude to accelerate pretraining research, which helps the company's models acquire their core knowledge and capabilities. It's the latest high-profile hire for Anthropic, which is poised to surpass OpenAI's private market valuation and is in an intensifying battle for talent with its chief AI rival. Ross Nordeen, a founding member of xAI and an ex-Tesla employee, announced earlier this month he was joining Anthropic, the same day the company struck a deal with Elon Musk's SpaceX to rent compute capacity at xAI's Colossus 1 data center in Memphis, Tennessee. After helping to start OpenAI, Karpathy left for Tesla in 2017 to serve as director of AI. There, he led the computer vision team for Tesla Autopilot. Musk recruited Karpathy away from OpenAI while the Tesla CEO was a board member at both tech companies. Karpathy's work at OpenAI and Tesla came up repeatedly during the Musk v. Altman trial, which concluded on Monday, with the jury and judge ruling in OpenAI CEO Sam Altman's favor. In one email exchange that was presented as an exhibit during the proceedings, Musk described Karpathy as "arguably the #2 guy in the world in computer vision," behind Ilya Sutskever, another OpenAI co-founder. "The OpenAI guys are gonna want to kill me, but it had to be done...," Musk wrote, regarding his hiring of Karpathy. Karpathy was one of several OpenAI

employees Musk borrowed from OpenAI to do months of free work at Tesla , where the development of self-driving vehicles wasn't going as quickly as promised. Karpathy left Tesla in 2022, and the company still doesn't sell a vehicle that's safe to use without a human driver ready to steer or brake at all times. After leaving Tesla, Karpathy briefly went back to OpenAI before starting AI education startup Eureka Labs, where he has worked until now. Karpathy holds a PhD in computer science from Stanford. WATCH: 2026 CNBC Disruptor 50: Watch the full interview with Anthropic co-founder Daniela Amodei

8. Andrej Karpathy's move to Anthropic is a major win for the AI giant in the talent wars - Business Insider

Business Insider | 2026-05-19

Andrej Karpathy, the OpenAI founding member who was formerly Tesla's AI boss, announced on Tuesday that he had joined Anthropic.

9. Influential AI researcher Andrej Karpathy announces he's joining Anthropic - Venturebeat

Venturebeat | 2026-05-19

Andrej Karpathy, the influential 39-year-old Slovak-Canadian AI researcher and one of the original 11 co-founders of OpenAI , and former head of Tesla's AI division, announced on Tuesday, May 19 that he's joining rival lab Anthropic. As Karpathy posted from his account on the social network X : "Personal update: I've joined Anthropic. I think the next few years at the frontier of LLMs will be especially formative. I am very excited to join the team here and get back to R&D. I remain deeply passionate about education and plan to resume my work on it in time." Anthropic's current Head of Pretraining, Nicholas Joseph, also a former OpenAI alumnus, added more context to Karpathy's new role at Anthropic in a post of his own on X, writing: " Excited to welcome Andrej to the Pretraining team! He'll be building a team focused on using Claude to accelerate pretraining research itself. I can't think of anyone better suited to do it — looking forward to what we build together!" An Anthropic spokesperson confirmed to VentureBeat via email that Karpathy will be starting a team focused on using Claude, Anthropic's own, increasingly popular AI model, to accelerate pretraining research. This would put Anthropic further toward the overarching AI research goal of many around the world to develop " recursive self-improvement, " that is, AI that is capable of training its successors or upgrading itself with increasingly lesser, or ultimately no human intervention. The announcement came on the same day as the start of rival AI-focused tech firm Google's annual I/O developer conference in its headquarters city of Mountain View, California, when many new releases and announcements were expected. Karpathy is widely known for spanning three parts of the modern AI boom: academic research, big-company deployment and online education. His own website describes him as an AI researcher and educator who was a founding member of OpenAI, later served as Director of AI at Tesla, and helped create Stanford's first deep learning course, CS231n. OpenAI's December 2015 launch announcement also listed Karpathy among the group's founding members. At Tesla, where he worked from 2017 to 2022, Karpathy led the computer vision team for Autopilot and says his team handled in-house data labeling, neural network training and deployment on Tesla's custom inference chip. He then returned to OpenAI from 2023 to 2024, where his website says he built a team focused on midtraining and synthetic data generation — experience directly relevant to Anthropic's reported pretraining role. Karpathy's academic work began at Stanford, where he earned his PhD under Fei-Fei Li and focused on neural networks for computer vision, natural language processing and the intersection of the two. He also interned at Google Brain, Google Research and DeepMind, according to his website. His education includes an

MSc from the University of British Columbia and a BSc from the University of Toronto, where he double-majored in computer science and physics. Since leaving OpenAI in 2024, Karpathy has become one of AI's most visible public educators, publishing technical and general-audience videos on large language models and neural networks. He also launched Eureka Labs in July 2024 as an "AI-native" school; its first product, LLM101n , is described as an undergraduate-level course guiding students through training their own AI system. Acting on his own as a free agent over the last two years, Karpathy has also helped push open source AI research forward with products and standards including autoresearch , an LLM-driven automated researcher that can run multiple hypothesis and experiments simultaneously, and the LLM Knowledge Base , an autonomous system of storing memory and context for AI agents in a kind of ever-growing library designed for them to access. The big question is what becomes of these and Karpathy's open source AI efforts more generally as he joins Anthropic, a lab that has supported open source via the launch of its Model Context Protocol (MCP) technical standard, but which also famously has shipped primarily proprietary AI models and harnesses (such as Claude and Claude Code). Based on the last statement in his announcement post on X — " I remain deeply passionate about education and plan to resume my work on it in time" — it appears that at least his contributions to the AI-native school effort will be paused as he digs in at Anthropic.

10. Andrej Karpathy, Tesla Alum and OpenAI Co-Founder, Joins Anthropic - WSJ

WSJ | 2026-05-19

Andrej Karpathy, Tesla Alum and OpenAI Co-Founder, Joins Anthropic WSJ

Analysis

Metric	Value
articles	0
news	10
images	1
keywords	20
facts	40
characters	24776

Top Keywords

- karpathy
- anthropic
- openai
- team
- tesla
- where
- andrej
- what
- coding
- will
- research
- vibe
- code
- model
- agent
- year
- work
- training
- education
- pretraining

References

- 1 https://en.wikipedia.org/wiki/Andrej_Karpathy
- 2 <https://www.forbes.com/sites/jodiecook/2026/06/12/is-vibe-coding-already-dead-even-karpathy-is-moving-on/>
- 3 <https://techcrunch.com/2026/05/19/openai-co-founder-andrej-karpathy-joins-anthropics-pre-training-team/>
- 4 <https://fortune.com/2026/05/19/who-is-andrej-karpathy-vibe-coding-anthropic-openai-rubiks-cube/>
- 5 <https://www.bloomberg.com/news/articles/2026-05-19/openai-founding-member-andrej-karpathy-takes-role-at-anthropic>
- 6 <https://www.reuters.com/business/autos-transportation/former-tesla-ai-executive-openai-founding-member-andrej-karpathy-joins-anthropic-2026-05-19/>
- 7 <https://www.axios.com/2026/05/19/anthropic-openai-karpathy-andrej-claude>
- 8 <https://www.cnbc.com/2026/05/19/anthropic-hires-openai-cofounder-andrej-karpathy-former-tesla-ai-lead.html>
- 9 <https://www.businessinsider.com/anthropic-hires-andrej-karpathy-2026-5>
- 10 <https://venturebeat.com/technology/andrej-karpathy-announces-hes-joining-anthropic>
- 11 <https://www.wsj.com/tech/ai/andrej-karpathy-tesla-alum-and-openai-co-founder-joins-anthropic-c665f51f>