

Open Source AI vs Closed AI

A Comprehensive Research Biography

11 July 2026

Contents

1. Introduction
2. Image Gallery
3. Career and Contributions
4. Recent Developments
5. References

Chapter 1

Introduction

Chapter 3

Career and Contributions

Chapter 4

Recent Developments and Current Work

AI: Alibaba open vs closed AI 'canary in the coalmine'. AI-RTZ #1143 - AI: Reset to Zero

AI: Reset to Zero | 2026-07-10

I've discussed for a long while how China's tech/AI companies large and small have taken the lead globally in open source AI models this AI Tech Wave . The lead was ceded by Meta's global Llama models lead a couple of years ago, as it shifted its focus to more closed efforts like Muse Spark models released recently. And China's DeepSeek and other AI companies have pressed the pedal to whatever AI metal they could get their hands on, to stay close to performance vs the leading frontier AI models from Anthropic and OpenAI in particular. And US companies have also been leaning into the open source AI models from China. Despite ongoing US/China AI machinations. The leader in open source AI from China of late has been Alibaba with its Qwen open source models. And their strategy between open and closed has been shifting. Echoing some of the moves by Meta in the US in particular. This is a shift worth noting, as Alibaba's challenges are a 'canary in the coal mine' type of moment for the world's leading AI companies, be they in China or the US . The NY Times lays it out well in "Alibaba's AI is a hit, but hard to turn into a moneymaker": "Alibaba became one of China's most valuable companies on the strength of its globe-spanning e-commerce business. But in recent years, it has been outmaneuvered by lower-cost rivals at home and has struggled to expand overseas. After Mr. Ma faded from public view, a new generation of executives took over, adopting a lower profile while emphasizing Alibaba's role within Beijing's broader technology policies." "In hindsight, Alibaba had gradually laid the groundwork to become an A.I. heavyweight. As its online shopping and logistics businesses expanded, the company built data centers capable of processing vast volumes of customer data worldwide." Alibaba over the last 25 plus years can be seen as an Amazon of China, with more of a B2B ecommerce focus in its early decades. But like Amazon with AWS on the cloud side, it has also expanded on that opportunity at stake. Additionally, Alibaba has been a fast follower to DeepSeek's world-shaking open source model debut in January of 2025, and echoed Meta's ramp a couple of years ago with open source Llama AI models. "It also borrowed a page from Amazon, parlaying its e-commerce success and technological infrastructure into a major cloud computing business, giving Alibaba two essential ingredients for building artificial intelligence systems: data and computing power." "Alibaba introduced its family of A.I. models, Qwen, in 2023 and quickly made them open source. When DeepSeek, an A.I. start-up, burst into the spotlight with claims that it had built a powerful model for a small fraction of the cost of Western competitors, global investors rushed to capitalize on the excitement around China's open-source A.I." "China has made open-source technology a pillar of its drive to become an A.I.

superpower . Nearly all of its leading A.I. systems are open source.” “By January, Qwen had become the world’s most downloaded open-source A.I. system. Its models were being downloaded about one million times a day, according to data from Hugging Face, which hosts many open-source A.I. projects.” “But that reach has not translated into big money. In the first three months of this year, Alibaba reported \$1.3 billion in revenue from A.I.-related products — less than 4 percent of its total revenue. That pales in comparison with the company’s plan to spend more than \$55 billion by the end of next year to build out its A.I. infrastructure.” “The push to turn its open-source success into a profitable business has quietly fractured the team behind Qwen.” That’s meant some hard internal decisions that led to the departure of key people who favored the open vs closed strategy. The one with less near-term AI monetization. “In March, Lin Junyang, Qwen’s lead engineer, announced that he was leaving Alibaba. Several other key engineers left around the same time. Two people familiar with the team said it had become divided by disagreements over how best to commercialize Qwen.” “Alibaba had long kept its largest, most advanced models proprietary, while releasing leading open-source models alongside them. Now, the company is signaling a broader shift away from widely used open-source models and toward closed ones that customers must pay to use. In April alone, Alibaba released three proprietary models within days of one another.” And they have their work cut out for them vs the US frontier AI companies. “The Qwen team remains focused on keeping pace with the leading models coming out of Silicon Valley, but there is a growing recognition that technological leadership will not be enough if the company doesn’t make money, said a member of Alibaba’s A.I. research lab, which includes the Qwen team and others, who spoke on the condition of anonymity because the person was not authorized to discuss internal matters.” “The challenge for Alibaba is apparent in its stock price. While A.I.-related stocks have surged in markets worldwide, Alibaba’s stock has fallen 37 percent this year in Hong Kong, where the broader market has declined 12 percent.” “Building a top-performing A.I. model is expensive, requiring enormous investments in hangar-size data centers filled with computer chips that consume copious amounts of electricity.” The other headwind of course is the US/China AI technology and trade curbs . “Chinese A.I. companies face an additional hurdle: U.S. export controls limit their access to the most advanced chips. China’s leading A.I. start-ups and researchers routinely say shortages of computing power are the biggest thing holding them back, and they are spending heavily to secure it.” “Those exorbitant costs are heaping financial pressure on Chinese A.I. companies. Start-ups like MiniMax and Z.ai have gone public to raise money from investors, and Alibaba is increasingly steering customers toward proprietary models even if that strategy upsets some of its top talent.” ““It is a tough choice to make, and one that every open-model lab has to make at some point,” said Kevin Xu, the founder of Interconnected Capital, a hedge fund that invests in A.I. technologies.” Alibaba in particular is a visible target for US attention at all levels. “Alibaba’s A.I. business faces mounting external challenges. The Pentagon recently placed the company on a blacklist of firms that it says support the Chinese military, a designation that Alibaba argues has already harmed its business. In a statement, Alibaba said it was not affiliated with the Chinese military.” And the US frontier AI companies have

also directed their ire at Alibaba and its China peers. “Anthropic and OpenAI have also accused Chinese companies, including Alibaba, of improperly harvesting data from their A.I. systems to accelerate the development of their models.” “Last month, Anthropic sent a letter to Senators Tim Scott, Republican of South Carolina, and Elizabeth Warren, Democrat of Massachusetts, accusing Alibaba of “brazenly” and “illicitly” trying to copy its technology using 24,000 fraudulent accounts. Alibaba declined to comment on the allegations.” The current monetization travails echo of course earlier tech cycles. I discussed this trend in AI Ramblings Daily (ARD) podcast #114 this week, titled “AI’s Pricing Umbrella getting larger”. “Many in the technology industry compare today’s A.I. boom to the early days of the dot-com era, when internet companies were flourishing but no one had yet figured out how to build durable businesses around the technology.” “Richard Lin, a vice president at the Silicon Valley company Datastrato who has long been involved in China’s open-source community, said China’s leading A.I. companies were all fighting for survival.” “There isn’t an A.I. company with a sustainable business model right now,” he said. “It’s not a healthy industry.” But the core takeaway is that Alibaba’s geopolitical path as a leading tech and AI company, with a global public ownership profile, is an important ‘canary in the coal mine’ for AI businesses open and closed. And we’re going to see similar examples on these dynamics this AI Tech Wave , both in China, the US and the rest of the world. Alibaba just happens to be one of the first going through this transition turbulence. Stay tuned. (NOTE: The discussions here are for information purposes only, and not meant as investment advice at any time. Thanks for joining us here)

Trump AI Restrictions Spur Shift Toward Open-Source Models - PYMNTS.com

PYMNTS.com | 2026-07-10

The Trump administration's recent restrictions on access to the most advanced U.S. artificial intelligence models are accelerating enterprise interest in open-source alternatives, including rapidly improving models developed in China. The trend underscores how government controls on frontier AI could reshape competition in the global AI market by encouraging businesses to diversify away from proprietary U.S. systems. A new report by Agence France-Presse (AFP), published in ET CIO, describes a sharp shift in industry sentiment following administration actions that limited access to leading closed AI models from Anthropic and OpenAI. While the White House has generally favored deregulation of the technology sector, the restrictions surprised many AI developers who had become accustomed to increasingly powerful models being released with relatively little government intervention. The developments have reignited debate over closed versus open AI models. Closed models, such as OpenAI's ChatGPT and Anthropic's Claude family, are controlled by their developers, who determine who may access them through subscription services or APIs. Open-source, or “open-weight,” models make their underlying model weights available for download, allowing organizations to run, modify and fine-tune them on

their own infrastructure without relying on continued access from the developer. Once released, those models cannot easily be withdrawn by either the vendor or governments. The administration's restrictions were particularly disruptive because they affected cutting-edge systems that many startups had integrated into their products. In early June, Anthropic was ordered to block non-U.S. users from accessing its flagship Mythos 5 and Fable 5 models. Faced with the operational complexity of verifying users' eligibility, the company instead withdrew the models entirely. Shortly afterward, OpenAI agreed to allow the federal government to approve customers seeking access to its newest GPT-5.6 model. The uncertainty has prompted companies to rethink their dependence on any single frontier model. Oren Michels, co-founder and chief executive of Barndoor AI, told AFP that businesses built exclusively around one proprietary model become significantly less reliable if access is suddenly interrupted. Haitham Mengad, co-founder of AI music startup Stems Labs, described Anthropic's Fable model as transformative for his company and said its withdrawal demonstrated the risks of relying too heavily on a single commercial provider. He later characterized the episode as a pivotal moment in convincing him that open-source alternatives deserved greater consideration. The report notes that open models were already gaining momentum because of rising costs associated with proprietary AI services. The administration's actions added another incentive by making access to leading commercial models appear less predictable. At roughly the same time, China's Zhipu AI (Z.ai) released GLM-5.2, an open model that reportedly approaches the performance of Anthropic and OpenAI's leading offerings on several industry benchmarks. AI analyst Andrew Curran told AFP the model's ability to be downloaded, customized and deployed on enterprise servers without licensing restrictions simultaneously lowers costs and reduces dependence on commercial frontier AI providers. Usage data cited in the report illustrates how rapidly the competitive landscape is changing. On OpenRouter, a platform that distributes requests across multiple AI models, the combined share of tokens consumed by Google, Anthropic and OpenAI fell from 55% in January to 33% by June. At the same time, non-U.S. models captured a growing portion of overall usage, with China's open-source DeepSeek emerging as the leading model on the platform by a substantial margin. The report also points to changing attitudes toward Chinese AI. While security concerns initially discouraged some companies from adopting Chinese-developed artificial intelligence models, several developers interviewed said those fears have diminished because organizations running open models locally retain full control over their data. Still, the report cautions that open-source AI may not remain insulated from government intervention forever. University of Pennsylvania Professor Ethan Mollick suggested that if frontier-level models are increasingly viewed as national security risks, governments beyond the United States could eventually seek to limit public release of their own most capable open models. "If Mythos-level models are considered risky, China will also not want them to be open," he said.

From Love Island to Precious Metals: How Prediction Markets Are Changing Finance
Cursor Prepares Workplace AI Agent to Challenge Claude
Cowork and ChatGPT Work Circle Secures National Trust Charter to Capture Digital Asset Custody Market

US crackdown on top AI fuels open-source surge - Tech Xplore

Tech Xplore | 2026-07-09

US crackdown on top AI fuels open-source surge Tech Xplore

Popular open source AI developer tool Ollama raises \$65M, grows to nearly 9M users - TechCrunch

TechCrunch | 2026-07-09

The popular open source AI tool Ollama has raised a \$65 million Series B, led by Theory Ventures, founder and CEO Jeff Morgan tells TechCrunch. This round follows a previous \$15 million Series A led by Benchmark's Peter Fenton. All told, the company has now raised \$88 million. Ollama, which launched in 2023, helps devs run open-weight AI models on their PCs, getting them up and running in minutes. It has been praised by developers across countless training sites, videos, blogs and social media posts. It has amassed 176,000 stars and nearly 17,000 forks on GitHub. Developers can also use Ollama to find models and access larger, more complex ones that it hosts on its neocloud via several subscription tiers, from free to \$100/month. It also tracks usage based on GPU time, not token limits. If the mission to help developers more easily build on their PCs sounds vaguely familiar, it should. Morgan and his co-founder Michael Chiang previously helped build Docker Desktop. They landed at Docker after it bought their previous startup, Kitematic. Docker makes containers that help cloud apps easy to move from cloud to cloud, or from desktop to cloud, abstracting away all the pesky hardware configuration issues. So Ollama essentially did for AI what Docker and Docker Desktop did for cloud. "Open models started coming out in 2023 but they were really hard to use," Morgan said. They had been geared toward researchers at the time, not programmers. "As a result, it was really hard to get them up and running." Three years after launching, Ollama is now "used by over 8.9 million developers every month, sitting in 85% of the Fortune 500 and growing like crazy," he said. All with only 14 employees. That career experience is what drew Benchmark's Peter Fenton to lead its earlier round and join the board. "What Jeff and Michael built with Docker is being used by 10 million-plus developers every day. The creative powers to create a product that goes to ubiquity for developers is extremely rare," Fenton told TechCrunch. Morgan and Fenton declined to discuss the startup's revenues and new valuation. However, Morgan says that the proving point for Ollama as a business happened around January, when OpenClaw became hot. That's when larger open models "suddenly became able to do these agentic tasks, like coding. Obviously, we saw the explosion of the assistants like OpenClaw, and this idea that open models can get real work done." Since then, the industry has been abuzz with the idea that paying users (particularly deep-pocketed enterprises and fast-growing AI application-layer startups) will increasingly turn to more affordable open models, reserving their use of closed models like Anthropic for more of an

as-needed basis. “I still think that this is the part that most of the debate gets wrong. It’s not an either/or,” Fenton says of open versus closed AI models. There will be plenty of business for both, he contends. However, every company with high inference expenses — the costs of using the models — has a “vital existential project” pushing them to move “to open-weight models,” he says. There’s plenty of evidence that such startups and enterprises are already turning to open models for their daily needs. That, obviously, bodes well for Ollama’s cloud business. But even more interesting, Ollama is another example of how AI is birthing a large new crop of open source projects that are turning into companies pursued by VCs. There are open source inference providers like Inferact, maker of vLLM, and RadixArk, maker of SGLang . There is OpenClaw and its alternatives like NanoClaw. There are even tiny startups building their own open models from scratch, like Arcee. To be sure, not every Ollama fan has been happy that the company has been pursuing making a living. About a year ago, a bunch of blog and social media posts complained that its cloud business was drawing attention away from its beloved free project and cited Ollama as an example of the so-called “Enshittification” of dev tools , as the trend is called. But Morgan sees its cloud service as an evolution of its open source mission to help programmers find and easily use models. Those state-of-the-art, large, open models are often “too big to run on your own computer. So we said, ‘Hey, let’s help find the compute for that,’” he explained. Board member Fenton adds, “Nothing has changed for the core product that’s free on the desktop. There’s zero change to the premise that this is the place you can discover and run local models.”

When you purchase through links in our articles, we may earn a small commission . This doesn’t affect our editorial independence. Last chance to save up to \$190 on TechCrunch Founder Summit. Join 1,000+ founders and VCs at all stages for real-world scaling insights and connections that move the needle. Savings end June 26, 11:59 p.m. PT .

Open-source AI models get a booster shot on US restrictions on frontier technology - ETEnterpriseai.com

ETEnterpriseai.com | 2026-07-09

Open-source AI: As US restrictions tighten on advanced AI models, open-source alternatives gain traction, driven by enterprises seeking cost-effective and reliable solutions for artificial intelligence.

AI open models have benefits. So why aren’t they more widely used? - MIT Sloan

MIT Sloan | 2026-01-20

When grocery shoppers find a generic product that’s 90% as good as the brand name version but costs 87% less, they usually put it in their carts. But when it comes to large language models, most artificial intelligence users pick the more expensive option. A new paper co-authored by Frank Nagle , a research scientist at the MIT Initiative on the Digital Economy , found that users

largely opt for closed, proprietary AI inference models, namely those from OpenAI, Anthropic, and Google. Those models account for nearly 80% of all AI tokens that are processed on OpenRouter, the leading AI inference platform. In comparison, less-expensive open models from the likes of Meta, DeepSeek, and Mistral account for only 20% of AI tokens processed. (A token is a unit of input or output to an AI model, roughly equivalent to one word in a prompt to an AI chatbot.) Open models achieve about 90% of the performance of closed models when they are released, but they can quickly close that gap — and the price of running inference is 87% less on open models. Nagle and co-author Daniel Yue at the Georgia Institute of Technology found that optimal reallocation of demand from closed to open models could cut average overall spending by more than 70%, saving the global AI economy about \$25 billion annually. “The difference between benchmarks is small enough that most organizations don’t need to be paying six times as much just to get that little bit of performance improvement,” Nagle said. “They need to think about how to use the right tool for the right job instead of defaulting to what’s popular.” Much time and effort goes into training frontier AI models, which are the models atop which inference models are built. There are architectures to design and trillion-token datasets to curate, and processing units powerful enough to operate continuously for months need to be acquired. To defray those costs, some companies keep their inference models proprietary, or closed. That means users must pay for the model and the underlying computing resources to access AI inference services. With a limited number of players in the market, high markups are all but inevitable. Open models make public one or more model details. Common examples include model weights, source code, training data, or architecture. That openness means that users can host and run models locally — even on a laptop, provided it’s powerful enough. “This openness enables inference at only the cost of compute power,” Nagle and Yue write, “essentially making the software free and creating competitive pressure akin to commodity markets.” Nagle likened the market of third parties building open AI inference models to the emergence of companies such as Red Hat, which offers software, training, and customer support atop the open-source Linux operating system. Users are able to reap the advantages of open source while vendors, with the advantage of economies of scale, assume the risk for issues like uptime and security. From May to September 2025, Nagle and Yue observed daily token usage on the OpenRouter site, which captures approximately 1% of all global spending on AI model inference. The researchers noted that the site is popular because users can engage with multiple providers of model inference through a single interface. “OpenRouter attracts the type of user who’s more likely to be willing to use open models, as the point is to be able to swap between models,” Nagle said. Even so, closed models still accounted for close to 80% of AI token usage over the five-month study period, as well as nearly 96% of the revenue that passed through OpenRouter. But closed models cost 87% more to run — \$1.86 per million tokens, on average, compared with 23 cents per million tokens for open models. Next, Nagle and Yue combined OpenRouter usage data with model performance benchmark data from Artificial Analysis and LM Arena. The former compares models against multiple performance metrics; the latter crowdsources user preferences. Their analysis showed that open models averaged 89.6% of closed-model performance but were usually able to close the gap within 13

weeks of a closed model's initial release. That figure had dropped from 27 weeks just one year prior. Using that information, the researchers calculated that, by switching to open alternatives superior to the closed models they are currently using, OpenRouter users could reduce costs more than 70% while improving benchmark performance by more than 14%. They extrapolated the impact for all AI inference models, using a market size estimate of more than \$35 billion from Menlo Ventures, determining that "optimal substitution to open models could save the AI industry approximately \$25 billion annually." If users could spend less on AI inference models and see better performance, why aren't they switching? Nagle said companies often have two sets of concerns. One set is valid; the other is based on misconceptions, he said. Among the valid concerns is that there may be significant costs associated with switching to a new model. These costs are not accounted for in the estimated \$25 billion in annual savings and differ greatly among users. "People build a whole ecosystem around an existing closed model, refine it, and build more on top of it. It's not going to be trivial to switch to an open model," Nagle said. "In the long run it may be cheaper, but in the short run it will be costly." Further, there may be reliability, regulatory, or security concerns that are easier to assuage with closed models than open models. For example, DeepSeek (an organization that produces open models) is based in China, and as a research organization, it's not subject to Chinese AI regulations for commercial products. The misconceptions are based on perceptions of inferior performance, which the researchers demonstrate are inaccurate, and a fear that using open models means private data suddenly becomes public. "That's not correct," Nagle said. "Open models can be built and run within your own infrastructure. Your data's never leaving your servers." Nagle encouraged companies to periodically review their use of AI inference models the same way they reevaluate software and infrastructure investments. If closed models are in place, it's worth considering whether there's a more cost-effective way to meet the company's needs. Open AI models matter for global economics and politics, not just corporate balance sheets, Nagle added. Nations that don't have a lot to spend will turn to frontier open and inference models. If these models come from China alone, then America's market influence may wane, as it has in Africa and Asia amid China's infrastructure investments. "The United States is investing in data centers but not data models," Nagle said. "It may benefit the U.S. to invest in creating frontier models to compete with China."

Frank Nagle is a research scientist at the MIT Initiative on the Digital Economy and the chief economist at the Linux Foundation. He studies how competitors can collaborate on the creation of core technologies, while still competing on products and services built on top of them — especially in the context of artificial intelligence. His work frequently explores the domains of open-source software, crowdsourcing, free digital goods, cybersecurity, and the generation of strategic predictions from unstructured big data. Daniel Yue is an assistant professor at the Georgia Institute of Technology. His research explores why firms openly share innovative knowledge without directly profiting from it, a strategy called "open disclosure."

Open source AI matters more than ever, says Hugging Face CEO Clem Delangue - CryptoRank

CryptoRank | 2026-07-10

Hugging Face CEO Clem Delangue says open-source AI is vital to avoid concentration of power, noting the platform hosts thousands of models and datasets and is used by roughly half of the Fortune 500 while the company turned down a large Nvidia investment last year to prioritize capital efficiency. He highlighted rising contributions from Chinese labs and argued that open-source transparency is critical for security and adoption in sensitive use cases like robotics, with implications for fundraising, token launch thinking and governance in crypto, DeFi, DEX and CEX contexts.

BitcoinWorld Open source AI matters more than ever, says Hugging Face CEO Clem Delangue Hugging Face CEO Clem Delangue believes the debate over open versus closed source AI has never been more important. In a recent episode of Bitcoin World's Equity podcast, Delangue argued that open source AI is not only thriving but essential for preventing a handful of large companies from controlling the future of the technology. Hugging Face has grown into a central hub for AI development, often described as a GitHub for machine learning. The platform hosts thousands of open models and datasets, and is now used by roughly half of the Fortune 500. Delangue noted that many companies begin their AI journey using proprietary APIs from major providers, but as they scale, the cost and flexibility advantages of open source models become increasingly attractive. The conversation comes at a time when the AI industry is grappling with questions of safety, transparency, and market concentration. Delangue pointed to Anthropic's recent decision to halt its Fable model release as an example of the tensions surrounding closed development. He warned that if only a few large corporations control the most advanced AI systems, the broader ecosystem could suffer from reduced innovation, higher costs, and less accountability. Open source models, by contrast, allow researchers, startups, and regulators to inspect and audit the technology. This transparency, Delangue argued, is critical for building trust and ensuring that AI benefits a wide range of users, not just the largest players. One of the more striking developments Delangue highlighted is the growing role of Chinese research labs in the open source AI space. A significant portion of the open models being downloaded in the United States now originate from Chinese institutions. Rather than viewing this as a reason to distrust open source, Delangue sees it as a challenge that the U.S. and its allies should address by investing more heavily in their own open source efforts. He emphasized that the solution is not to restrict open source, but to compete within it. Encouraging more domestic contributions to open model development could help maintain a balanced and diverse AI ecosystem. Delangue also discussed Hugging Face's approach to fundraising, which stands apart from the typical Silicon Valley playbook. The company turned down a large investment from Nvidia last year, opting instead to prioritize capital efficiency and long-term independence. Delangue explained that staying lean allows Hugging Face to focus on its mission of democratizing AI rather than chasing growth at any cost. While much of the public conversation around AI focuses on chatbots and

coding tools, Delangue believes robotics presents an even more urgent case for openness. A robot operating in a home or interacting with family members will have access to highly sensitive personal data. Without transparent, auditable AI systems, users have little way of knowing how that data is being used or what decisions the robot might make. Delangue argued that open source models can provide the scrutiny needed to ensure safety and privacy in robotics applications. As robots become more common in everyday life, the demand for open, accountable AI will only grow. The open source AI movement is gaining momentum, driven by cost pressures, transparency concerns, and the need for broader access to advanced technology. Hugging Face, under Delangue's leadership, continues to position itself at the center of this shift. The company's emphasis on capital efficiency, its willingness to address the dominance of Chinese models, and its focus on robotics all point to a future where open source plays a defining role in how AI is built and used. Q1: What is Hugging Face? Hugging Face is a platform that hosts open source AI models and datasets, similar to how GitHub hosts code. It is widely used by researchers and companies to share, download, and collaborate on machine learning projects. Q2: Why does Clem Delangue think open source AI matters more now? Delangue argues that without open source, a small number of large companies could control AI development, limiting innovation and transparency. Open source allows for broader access, lower costs, and public scrutiny of AI systems. Q3: How is Hugging Face different from other AI companies? Hugging Face prioritizes capital efficiency over rapid fundraising, turning down large investments to maintain independence. It also focuses on open source as a core principle, contrasting with companies that keep their models proprietary. This post Open source AI matters more than ever, says Hugging Face CEO Clem Delangue first appeared on BitcoinWorld .

Axios C-Suite: Open-source AI pits cost against security - Axios

Axios | 2026-06-22

Microsoft is considering using a Chinese AI model inside your office software — amplifying how your AI strategy is both a security and a cost calculation now. Why it matters: Unlike closed models such as Anthropic's Claude or OpenAI's ChatGPT, open-source AI is free to download, cheaper to run, and dominated by China . Between the lines: Microsoft exploring a secured version of China-built DeepSeek V4 to power Copilot Cowork — its agentic assistant and the most compute-hungry part of Microsoft 365 — shows that even the richest software company on Earth can't hold the line on closed American-only models. You probably can't either. The security catch: Microsoft's "secured" version would keep your data on its cloud platform Azure, not Chinese servers, which means the open-source question is shifting from if to how. "Open source" means the model weights — the actual trained intelligence — are public. Anyone can download, modify and deploy them. The best-performing, most-used ones are Chinese. The part you may not see: Many CEOs reading this are already running Chinese models somewhere in their stack and don't

know it, because their vendors and engineers chose them on price. Worth watching: Flo Crivello, CEO of AI agent platform Lindy, switched his company entirely to DeepSeek V4, citing millions in savings and better performance on core use cases — while keeping it on American soil, hosted by a U.S. provider rather than Chinese sources. The bottom line: Open-source Chinese models are good enough for most work at a fraction of the price. But you've got to be clear-eyed about their use. Why it matters: Your company likely already faces thorny China-centric compliance questions — especially if you've staffed your tech team with AI enablers. Between the lines: The fear isn't just that Beijing has a potential backdoor into your company. It's the regulatory threat to your entire workflow as U.S.-China competition escalates. We've been here before: Banning Huawei from the telecom networks of the U.S. and its closest allies cost billions to rip out and replace. Zoom out: This isn't a fringe risk. Chinese companies — DeepSeek, Xiaomi, MiniMax, Tencent, Alibaba — now dominate OpenRouter, the industry's most honest real-time scoreboard of developer usage, with DeepSeek at No. 1. Worth noting: Not all open models are Chinese. Meta's Llama and Europe's Mistral are credible alternatives with no ties to Beijing. America made the opposite bet of China: We went all-in on our ahead-of-the-pack frontier proprietary models. The big picture: The U.S. incentivized its frontier AI labs to optimize for revenue — with eyes on \$1T+ IPOs. That pushed them toward subscriptions and closed ecosystems where every use can be metered. Meta was the major open exception on American shores. Meta's 2023 release of Llama 2 was a breakthrough, and Nvidia CEO Jensen Huang later called it "probably the biggest event in AI" from that year. The bottom line: In the U.S., economics mean the AI feedback loop has to be mostly proprietary: Each lab learns from its own users, then sells the improvements back to them. Open-source advocates think American frontier labs are begging for regulation in order to corner the AI market. Their claim: Costly compliance requirements — think: safety audits or liability frameworks — would be trivial for companies like Anthropic or OpenAI and fatal to any open-source project. What they're saying: Former White House AI czar David Sacks, who still has President Trump's ear on cutting-edge issues, outlined this "regulatory capture agenda" on last weekend's episode of the "All-In" podcast. Yes, but: Sacks was bemoaning the guardrails Anthropic had established around discussing cybersecurity and biology with its powerful Fable model. But the episode was recorded before the Trump administration issued export controls that caused Anthropic to shut down consumer access to it. The bottom line: For the moment, the frontier labs are feeling the sting of government regulation long before any open-source model creators. I asked Misha Laskin, co-founder and CEO of Reflection AI, an American open frontier lab, to make the pro-open-source case to other CEOs. Open models are a geopolitical priority. They're Trojan horses for the infrastructure they bring with them. When a nation or enterprise adopts an open model, it also adopts the chips and software stack that come with it natively. ■ If you're a CEO or on a CEO's team: Ask to join Jim's new weekly Axios C-Suite newsletter.

Chinese A.I. Models Close the Gap With Anthropic and OpenAI - The New York Times

The New York Times | 2026-07-01

Chinese A.I. Models Close the Gap With Anthropic and OpenAI The New York Times

A new, inexpensive Chinese AI model is catching up with Anthropic, OpenAI on their home turf - Reuters

Reuters | 2026-07-02

A new, inexpensive Chinese AI model is catching up with Anthropic, OpenAI on their home turf
Reuters

References

- 1 <https://michaelparekh.substack.com/p/ai-alibaba-open-vs-closed-ai-canary>
- 2 <https://www.pymnts.com/news/artificial-intelligence/2026/trump-ai-restrictions-spur-shift-toward-open-source-models/>
- 3 <https://techxplore.com/news/2026-07-crackdown-ai-fuels-source-surge.html>
- 4 <https://techcrunch.com/2026/07/09/popular-open-source-ai-developer-tool-ollama-raises-65m-grows-to-nearly-9m-users/>
- 5 <https://enterpriseai.economictimes.indiatimes.com/news/industry/open-source-ai-models-surge-amidst-us-restrictions-on-advanced-technologies/132283304>
- 6 <https://mitsloan.mit.edu/ideas-made-to-matter/ai-open-models-have-benefits-so-why-arent-they-more-widely-used>
- 7 <https://cryptorank.io/news/feed/23747-open-source-ai-matters-hugging-face-ceo-clem-delange>
- 8 <https://www.axios.com/2026/06/22/open-source-ai-china-cost-risk-glm-deepseek>
- 9 <https://www.nytimes.com/2026/06/25/technology/zai-china-artificial-intelligence-models.html>
- 10 <https://www.reuters.com/world/china/a-new-inexpensive-chinese-ai-model-is-catching-up-with-anthropic-openai-their-2026-07-02/>